



Object Detection with Vertex AI and AutoML

Daniel Bank



<https://github.com/danielbank/object-detection>



Scope of This Talk

- Brief history of object detection solutions
- Build an object detection model using Vertex AI
- Run object detection on an Android device (no network calls)
- A little TensorFlow Lite along the way

TL;DR:

It is incredibly easy* to build an object detection model with Vertex AI and deploy it on Android

* assuming you are willing to spend time labeling your data and money training a model

What is Object Detection?



☐ **Image classification (Single-label)**

Predict the one correct label that you want assigned to an image.



☐ **Image classification (Multi-label)**

Predict all the correct labels that you want assigned to an image.





Image object detection

Predict all the locations of objects that you're interested in.




☐ **Image segmentation**

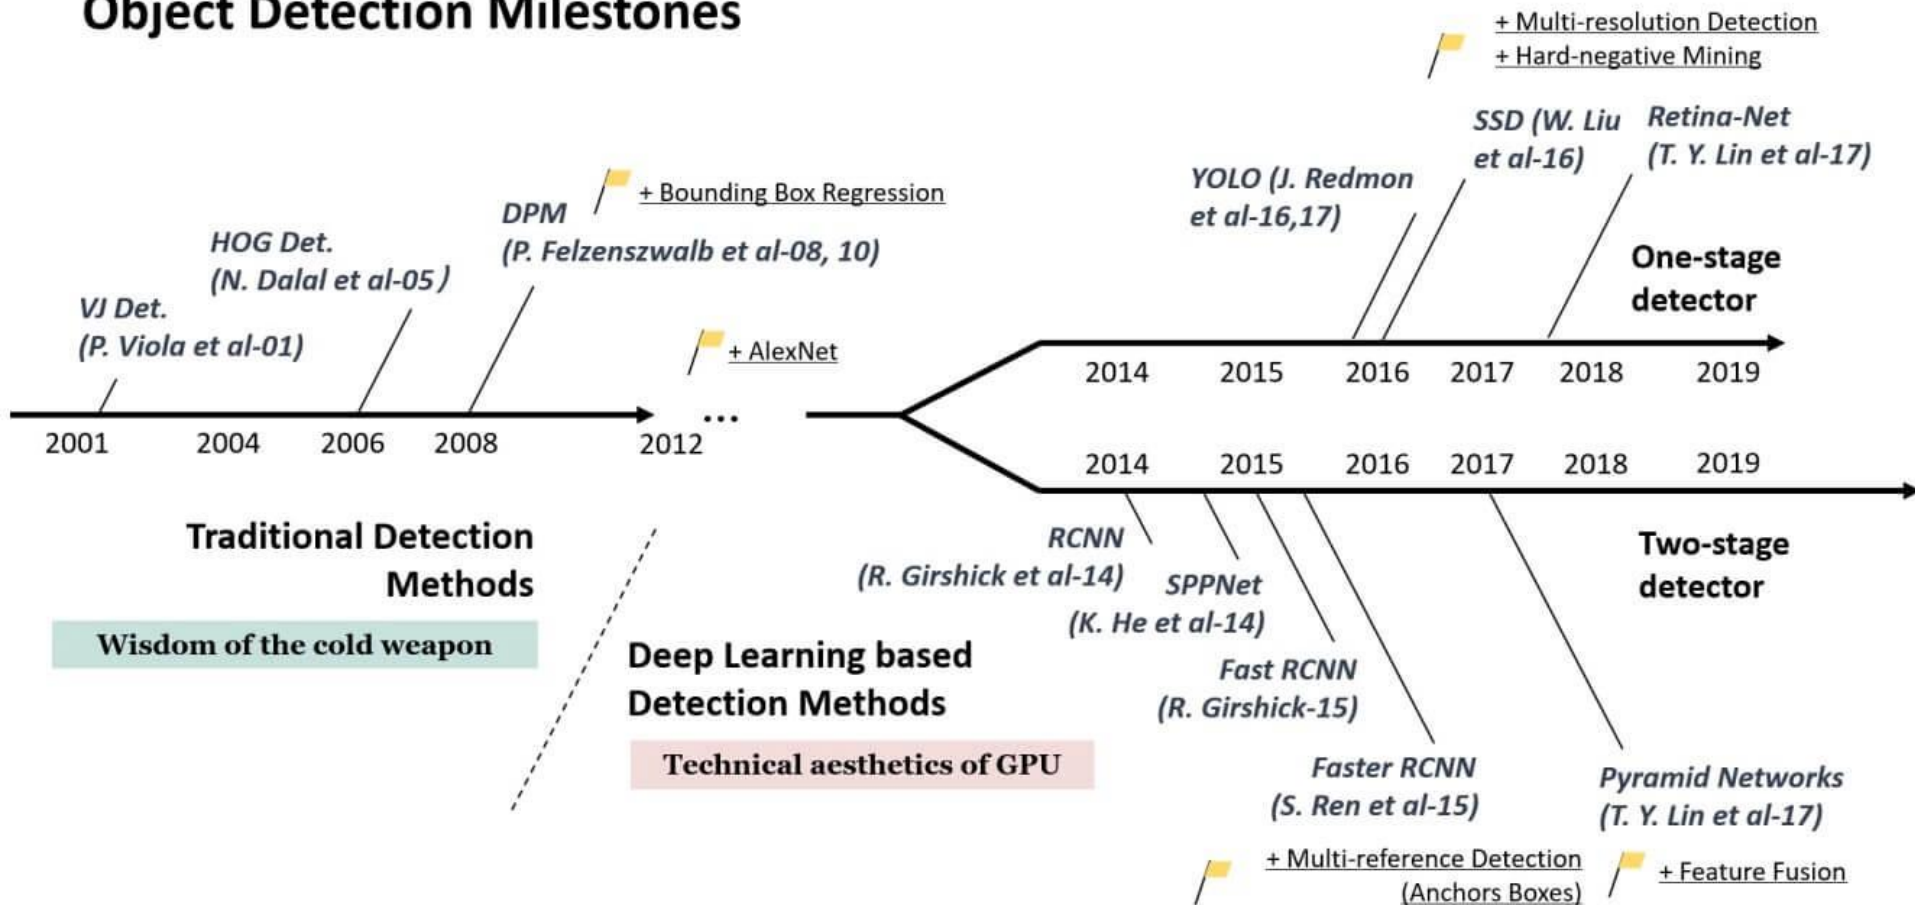
Predict per-pixel areas of an image with a label.

OpenCV

- C++ Library for computer vision
- Initially released in 2000
- Image processing, video capture, and analysis
- DNN Module (Deep Neural Network) in 2015



Object Detection Milestones



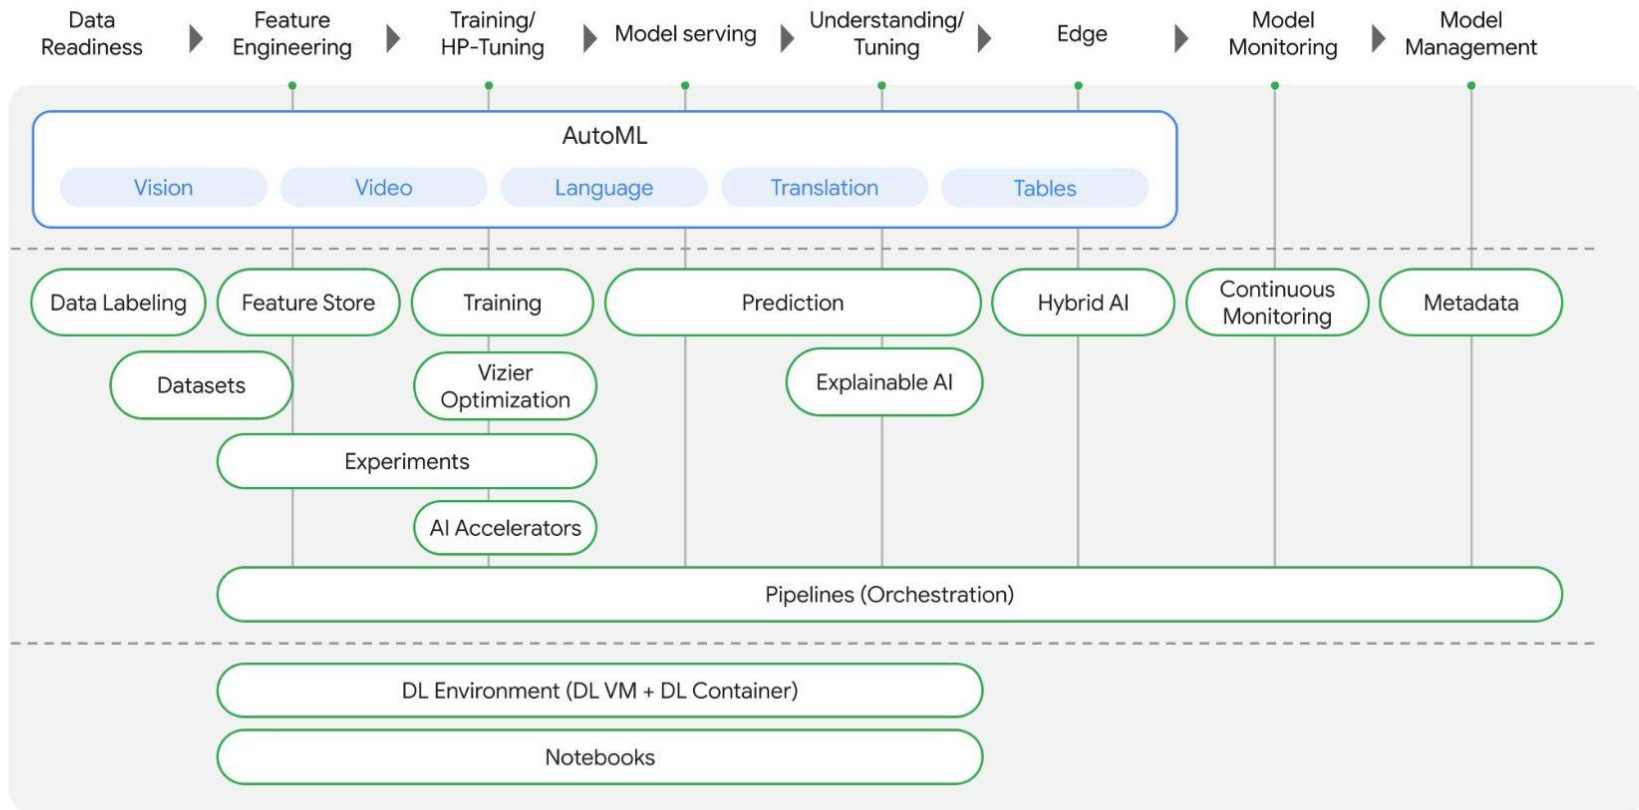
TensorFlow Object Detection API

- An open source framework built on top of TensorFlow that makes it easy to construct, train and deploy object detection models
- Released in 2017
- Includes a [Detection Model Zoo](#), a collection of pre-trained models
 - Model Zoo only has two-stage detectors

Vertex AI and AutoML

- Vertex AI
 - Google's Managed ML Platform
 - Launched in 2021
- Neural Architecture Search (NAS) a.k.a AutoML
 - AI that generates other Neural Networks
 - Developed by Google Brain in 2017
 - Vertex AI NAS
 - Available for Qualcomm Technologies Neural Processing SDK, optimized for Snapdragon 8

What's included in Vertex AI?



Building a Dataset

Detecting Dice

- Six-sided Dice Dataset: <https://www.kaggle.com/nellbyler/d6-dice>
 - 250 images
 - 1-25 six-sided dice per image
 - Images and annotations (YOLO format)
- Dice Detection Tutorial: https://github.com/nell-byler/dice_detection



Vertex AI



Dashboard



Datasets



Feature Store



Labeling tasks



Workbench



Pipelines



Training



Experiments



Model Registry



Endpoints



Batch predictions



Metadata



Matching Engine



Marketplace




Create dataset

☐ Image classification (Single-label)


Predict the one correct label that you want assigned to an image.

☐ Image classification (Multi-label)

Predict all the correct labels that you want assigned to an image.

☒ Image object detection

Predict all the locations of objects that you're interested in.

☐ Image segmentation

Predict per-pixel areas of an image with a label.

Region

us-central1 (Iowa) ▾ ?

▼ ADVANCED OPTIONS

You can use this dataset for other image-based objectives later by creating an annotation set. [Learn more](#)



IMPORT

BROWSE

ANALYZE


All	250
Labeled	0
Unlabeled	250

Filter Filter labels +


Images ▾

ADD NEW LABEL

Filter Filter items

☐ Select all


No bounding boxes



No bounding boxes



No bounding boxes



No bounding boxes

Training jobs and models
















Use this dataset and annotation set to train a new machine learning model with AutoML or custom code


TRAIN NEW MODEL

Labeling tasks

If your data still needs to be labeled, create a labeling task to have others label it for you

CREATE LABELING TASK


-  Vertex AI
-  Dashboard
-  Datasets
-  Feature Store
-  Labeling tasks
-  Workbench
-  Pipelines
-  Training
-  Experiments
-  Model Registry
-  Endpoints
-  Batch predictions
-  Metadata
-  Matching Engine
-  Marketplace

 Item 5 of many


Data split
Default


OBJECTS

DETAILS


 Filter

Filter labels




 1 (1)


1 1

 4 (1)


4 1

 2 (1)

2 1

 6 (3)



6 1
6 2
6 3

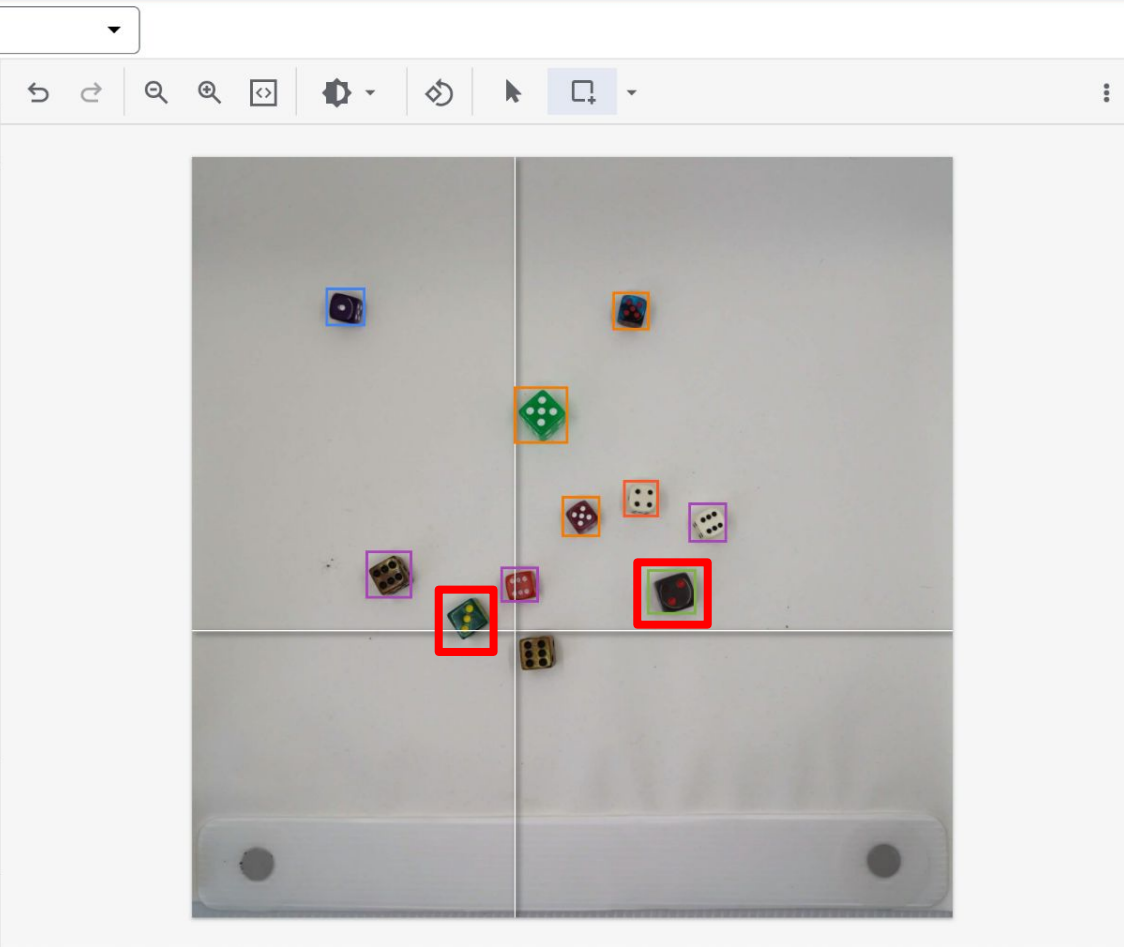
 5 (3)

5 1
5 2

ADD LABEL

SAVE



Vertex AI

Dashboard

Datasets

Feature Store

Labeling tasks

Workbench

Pipelines

Training

Experiments

Model Registry

Endpoints

Batch predictions

Metadata

Matching Engine

Marketplace

Item 5 of many

Data split
Default ▾

OBJECTS

DETAILS

Filter

Filter labels

+

1 (1)

1 1

4 (1)

4 1

2 (1)

2 1

6 (4)

6 1

6 2

6 3

6 4

5 (3)

5 1

ADD LABEL

SAVE

<

>

↶

↷

🔍

🔍

📄

⚙️ ▾

🔄

🖱️

📐 ▾

⋮

Filter |Type to filter

1

4

2

6

3

5

Add label...

Training a Model



Vertex AI



IMP

All

Label

Unlab

F

Imag

1

2

3

4

5

6

ADI

TOOLS



Dashboard



Workbench



Pipelines

DATA



Feature Store



Datasets



Labeling tasks

MODEL DEVELOPMENT

DEPLOY AND USE



Marketplace

<|

Train new model

1 Training method

2 Model details

3 Training options

4 Compute and pricing

START TRAINING

CANCEL

Dataset

d6_dice_1662878099415



Annotation set

d6_dice_1662878099415_iod



Objective

Image object detection



Please refer to the pricing guide for more details (and available deployment options) for each method.

Model training method

☐ AutoML

Train high-quality models with minimal effort and machine learning expertise. Just specify how long you want to train. [Learn more](#)

☒ AutoML Edge

Train a model that can be exported for on-prem/on-device use. Typically has lower accuracy. [Learn more](#)

☐ Custom training (advanced)

Run your TensorFlow, scikit-learn, and XGBoost training applications in the cloud. Train with one of Google Cloud's pre-built containers or use your own. [Learn more](#)

CONTINUE



IMP

All

Label

Unlab

F

Image

1

2

3

4

5

6

ADI

TOOLS



Dashboard



Workbench



Pipelines

DATA



Feature Store



Datasets



Labeling tasks

MODEL DEVELOPMENT

DEPLOY AND USE



Marketplace



Train new model

☒ Training method☒ 2 Model details☐ 3 Training options☐ 4 Compute and pricing

START TRAINING

CANCEL

☒ Train new model

Creates a new model group and assigns the trained model as version 1

☐ Train new version

Trains model as a version of an existing model

Name *

d6_dice_edge

Description

D6 Dice Model for Edge Devices

Data split

☒ Randomly assigned☐ Manual (Advanced)

Your dataset will be automatically randomized and split into training, validation, and test sets using the following ratios. [Learn more](#)

Training

80

%

Validation

10

%

Test

10

%

☒ Training: 80%☐ Validation: 10%☐ Test: 10%☐ Default: 0%

Encryption

☒ Google-managed encryption key

No configuration required

☐ Customer-managed encryption key (CMEK)

Manage via Google Cloud Key Management Service

[SHOW LESS](#)[CONTINUE](#)

TOOLS

Dashboard

Workbench

Pipelines

DATA

Feature Store

Datasets

Labeling tasks

MODEL DEVELOPMENT

DEPLOY AND USE

Marketplace

<|

Train new model

✓ Training method

✓ Model details

3 Training options

4 Compute and pricing

START TRAINING

CANCEL

	Goal	Package size	Accuracy	Latency on iPhone X ▾
<input type="radio"/>	Higher accuracy	5.6 MB	Higher	34ms
<input checked="" type="radio"/>	Best trade-off	3.1 MB	Medium	23ms
<input type="radio"/>	Faster predictions	557 KB	Lower	8ms

Please note that prediction latency estimates are for guidance only. Actual latency depends on your network connectivity. Edge TPU predictions typically will have lower latency.

CONTINUE

TOOLS

Dashboard

Workbench

Pipelines

DATA

Feature Store

Datasets

Labeling tasks

MODEL DEVELOPMENT

DEPLOY AND USE

Marketplace

<|

Train new model

✓ Training method

✓ Model details

✓ Training options

4 Compute and pricing

START TRAINING

CANCEL

Enter the **maximum** number of node hours you want to spend training your model.

You can train for as little as 1 node hour. You may also be eligible to train with free node hours. [Pricing guide](#)

Budget * 2 Maximum node hours

Estimated completion date: Sep 18, 2022 4 PM GMT-7

☒ Enable early stopping

Ends model training when no more improvements can be made and refunds leftover training budget. If early stopping is disabled, training continues until the budget is exhausted.

Google Cloud

Histogramo

Search Products, resources, docs (/)

1

?

Vertex AI

TOOLS

Dashboard

Workbench

Pipelines

DATA

MODEL DEVELOPMENT

Training

Experiments

Metadata

DEPLOY AND USE

Endpoints

Model Registry

Batch predictions

Matching Engine

Marketplace

Training

+ CREATE

REFRESH

LEARN

TRAINING PIPELINES

CUSTOM JOBS

HYPERPARAMETER TUNING JOBS

Training pipelines are the primary model training workflow in Vertex AI. You can use training pipelines to create an AutoML-trained model or a custom-trained model. For custom-trained models, training pipelines orchestrate custom training jobs and hyperparameter tuning with additional steps like adding a dataset or uploading the model to Vertex AI for prediction serving. [Learn More](#)

Region
us-central1 (Iowa)

Filter Enter a property name

Name	ID	Status	Job type	Model type	Created	Elapsed time	Labels
d6_dice_edge	3563755504767336448	<div>Training</div>	Training pipeline	Image object detection	Sep 18, 2022, 1:38:12 PM	28 min 34 sec	—

Google Cloud

Histogramo

Search

Products, resources, docs (/)

1

?

Vertex AI

TOOLS

Dashboard

Workbench

Pipelines

DATA

MODEL DEVELOPMENT

Training

Experiments

Metadata

DEPLOY AND USE

Marketplace

Training

+ CREATE

REFRESH

LEARN

TRAINING PIPELINES

CUSTOM JOBS

HYPERPARAMETER TUNING JOBS

Training pipelines are the primary model training workflow in Vertex AI. You can use training pipelines to create an AutoML-trained model or a custom-trained model. For custom-trained models, training pipelines orchestrate custom training jobs and hyperparameter tuning with additional steps like adding a dataset or uploading the model to Vertex AI for prediction serving. [Learn More](#)

Region

us-central1 (Iowa)


Filter


Enter a property name


Name	ID	Status	Job type	Model type	Created	Elapsed time	Labels
d6_dice_edge	3563755504767336448	Finished	Training pipeline	Image object detection	Sep 18, 2022, 1:38:12 PM	2 hr 18 min	—

Vertex AI

TOOLS


 Dashboard


 Workbench


 Pipelines

DATA


MODEL DEVELOPMENT


 Training


 Experiments


 Metadata


DEPLOY AND USE

 Endpoints

 Model Registry

 Batch predictions

 Matching Engine

 Marketplace

<1

←

d6_dice_edge

>

Version 1

VIEW DATASET

EXPORT

EVALUATE

DEPLOY & TEST

BATCH PREDICT

VERSION DETAILS

←

untitled_4497850143629901824

COMPARE

CREATE EVALUATION

Filter

Filter labels

All labels

0

6

0.773

1

0.765

3

0.746

2

0.693

5

0.677

4

0.623

Confidence threshold

?

0.34

IoU threshold

?

0.44

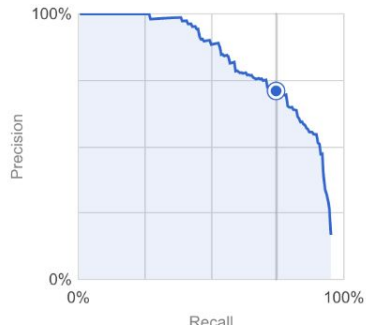
Test images

26

To evaluate your model, set the confidence threshold to see how precision and recall are affected. The best confidence threshold depends on your use case. Read some [example scenarios](#) to learn how evaluation metrics can be used.

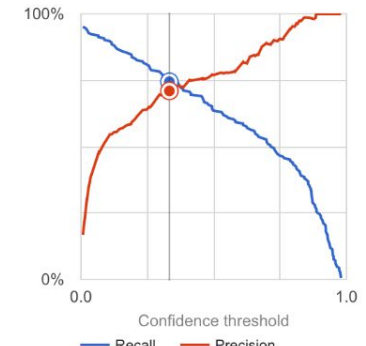
Precision-recall curve

?



Precision-recall by threshold

?



Google Cloud

Histogramo

Search

Products, resources, docs

Vertex AI

Tools

Dashboard

Workbench

Pipelines

Data

Model Development

Training

Experiments

Metadata

Deploy and Use

Endpoints

Model Registry

Batch predictions

Matching Engine

Marketplace

←

d6_dice_edge

Version 1

VIEW DATASET

Evaluate

Deploy & Test

Batch Predict

Version

TF Lite

Export your model as a TF Lite package to run your model on edge or mobile devices.

Container

Export your model as a TF S Model to run on a Docker c

Deploy your model

Endpoints are machine learning models made available for online prediction reques are useful for timely predictions from many users (for example, in response to an a request). You can also request batch predictions if you don't need immediate result

Deploy to Endpoint

Name	ID	Status	Models	Region	Monitoring
No active endpoints containing this model					

Test your model

PREVIEW

In order to test your model, you will need to deploy it first. [Pricing](#)

Export model

LEARN

The **TensorFlow Lite (.tflite)** format allows you to run your model on mobile and embedded devices.

- Export your model as a TF Lite package.

Destination folder on Cloud Storage *

gs:// d6_dice_edge

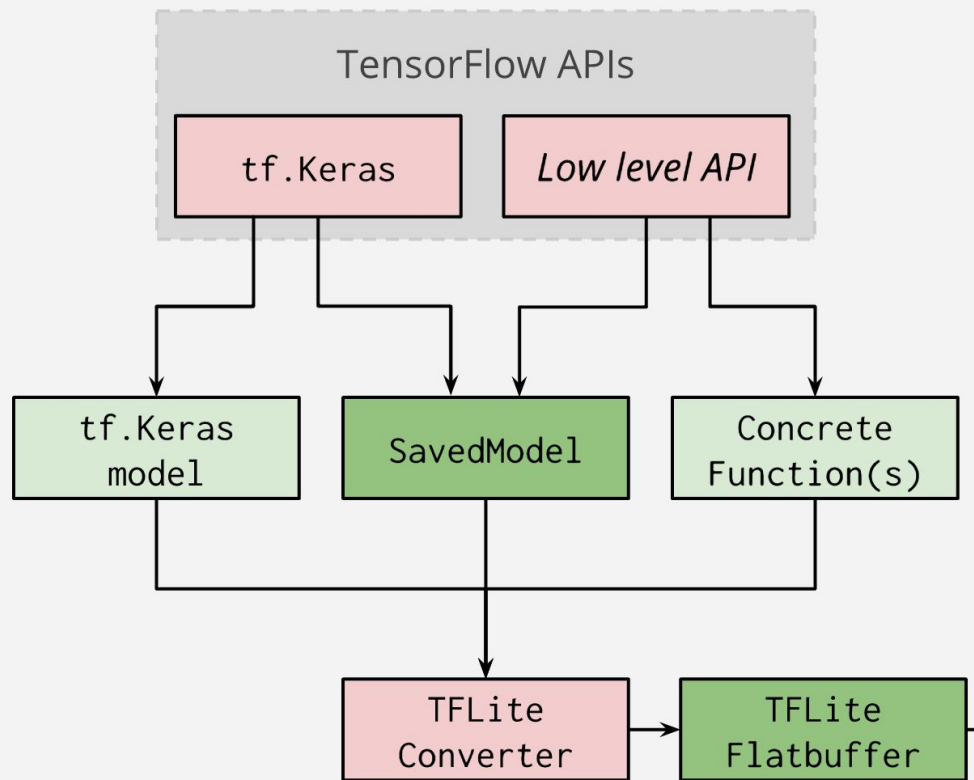
BROWSE

EXPORT

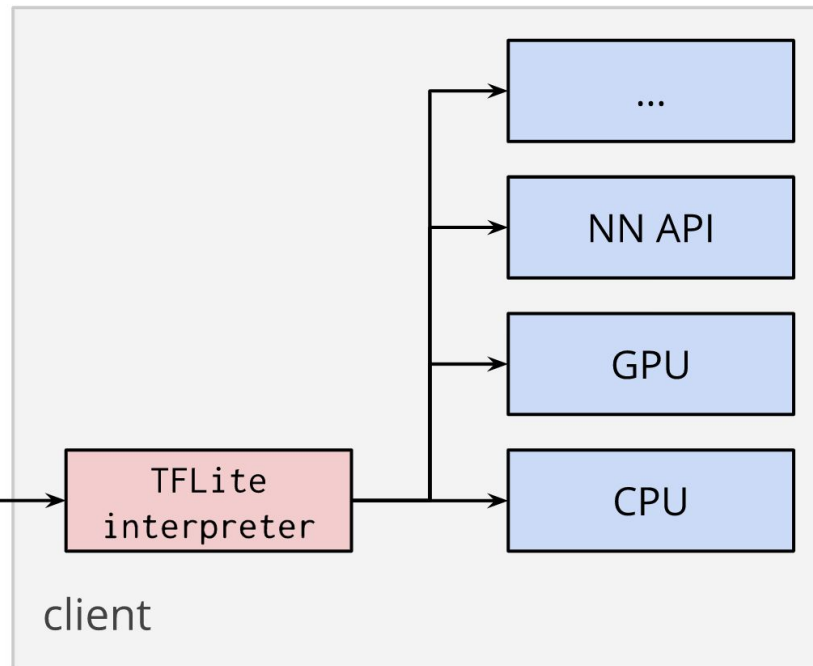

VIEW FOLDER
- Model export takes a couple of minutes. After exporting is finished, copy the package to your computer using the following command:

```
$ gsutil cp -r gs://d6_dice_edge ./download_dir
```

CLOSE



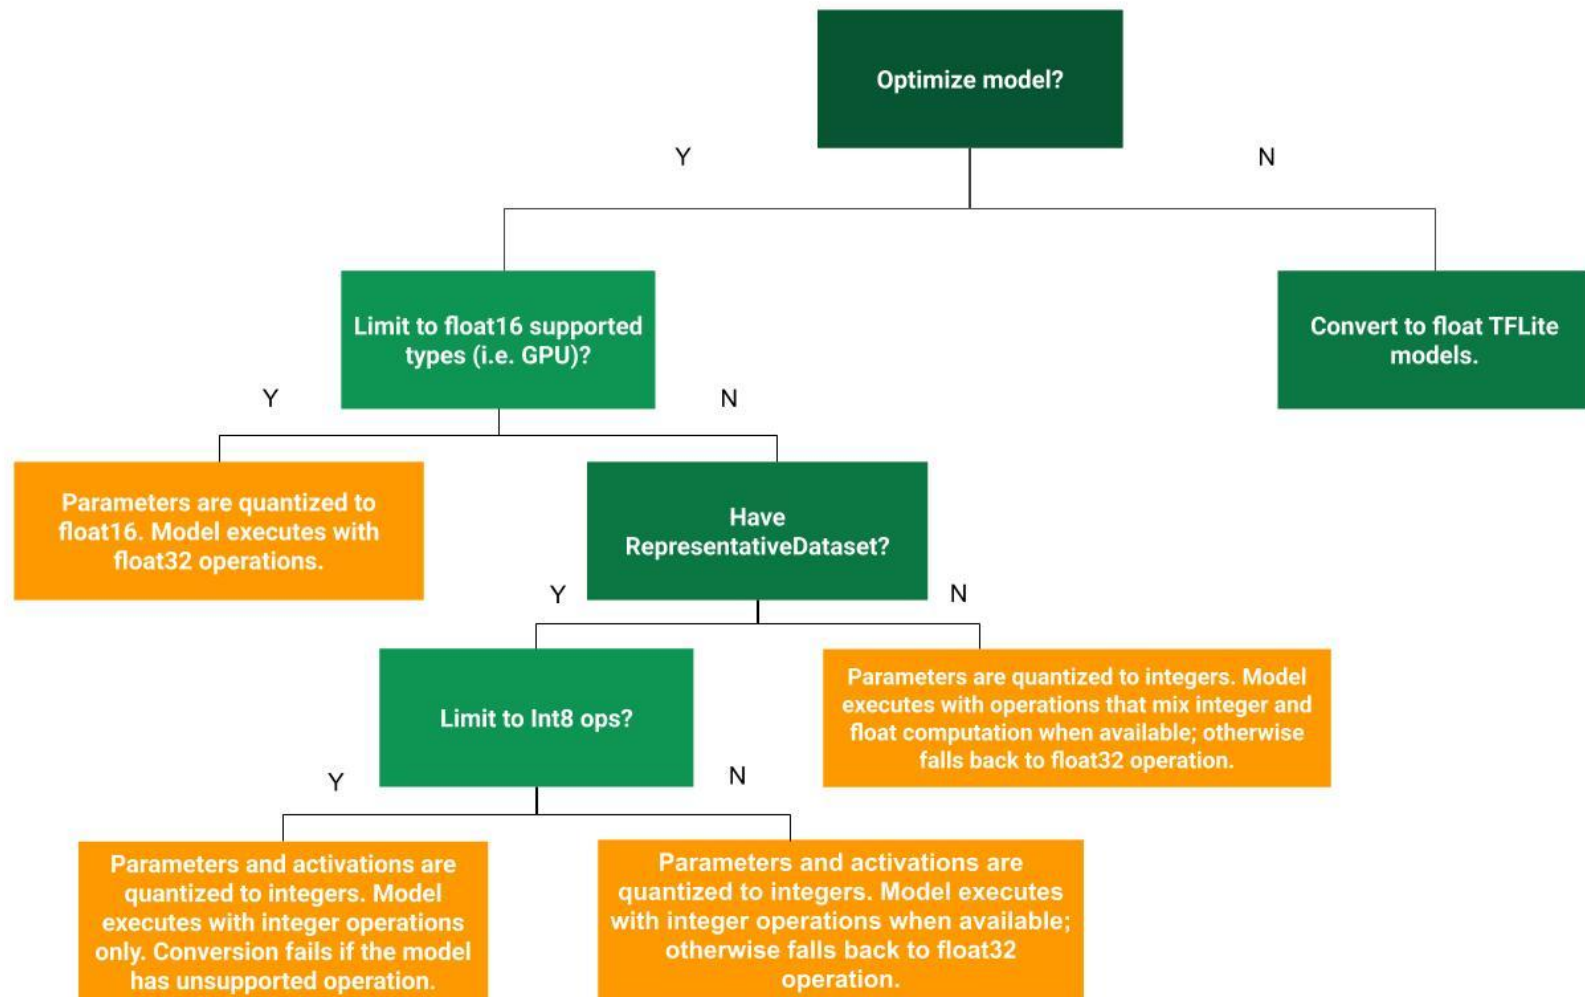
server



client

Post-Training Quantization Options

TF Lite Option	Technique Used	Benefits	Hardware
OPTIMIZE_FOR_SIZE	“Hybrid operations”	4x smaller, 2-3x speedup, accuracy	CPU
DEFAULT	Integer Quantization	4x smaller, More speedup	CPU, Edge TPU, etc.






THE LABELS ARE IN THE MODEL FILE

TFLite Android Example App

TensorFlow Lite Object Detection Android Demo

https://github.com/tensorflow/examples/tree/master/lite/examples/object_detection/android



```
+++ b/android/app/src/main/java/org/tensorflow/lite/examples/objectdetection/ObjectDetectorHelper.kt
@@ -82,14 +82,7 @@ class ObjectDetectorHelper(
```

```
    optionsBuilder.setBaseOptions(baseOptionsBuilder.build())
```

```
-    val modelName =
-        when (currentModel) {
-            MODEL_MOBILENETV1 -> "mobilenetv1.tflite"
-            MODEL_EFFICIENTDETV0 -> "efficientdet-lite0.tflite"
-            MODEL_EFFICIENTDETV1 -> "efficientdet-lite1.tflite"
-            MODEL_EFFICIENTDETV2 -> "efficientdet-lite2.tflite"
-            else -> "mobilenetv1.tflite"
-        }
+    val modelName = "model.tflite"
```






1 0.91



1 0.82



Links

- Object Detection
 - [Object Detection in 20 Years: A Survey](#)
 - [Object Detection using YOLOv5 and OpenCV DNN in C++ and Python](#)
 - [YOLO v5](#)
 - [YOLO Algorithm and YOLO Object Detection](#)
 - [COCO Dataset](#)
- [Vertex AI](#)
 - [Vertex AI NAS Announcement 11/30/2021](#)
 - [AutoML Beginner's Guide](#)
- TensorFlow and TensorFlow Lite
 - [TensorFlow Lite Object Detection Android Demo](#)
 - [Easier object detection on mobile with TensorFlow Lite](#)
 - [TensorFlow post-training quantization](#)
- Dice Detection
 - [Six-sided Dice Dataset](#)
 - [Real-time dice detection and classification](#)

Thank you!